

selective-attention decoder: 4 layers, width 512, 8 heads

